



GROUP 28 · CSIT321

System Architecture

Certacito.ai · A4 Prototype Presentation appendix

Project Certacito.ai, AI agent governance platform

Document architecture.md

Date 07 August 2026

Live system <http://20.92.93.30>

Certacito.ai is a governance layer that sits between an AI agent and the actions it wants to take. Every tool call is intercepted before execution, evaluated against policy that lives outside the model, and recorded in a tamper-evident log.

This describes what is actually built and deployed, not a target state.

The layers

```
Agent (OpenClaw)
| tool call proposed
v
Interception hook      certacito-hook / certacito-exec-gate
| POST /api/v1/intercept (X-API-Key)
v
Governance API          FastAPI, Python 3.13
|-- semantic guardrail  (runs first, can short-circuit)
|-- risk classifier
|-- policy engine        (fail-closed)
|-- approval service     (on ESCALATE)
v
Persistence             PostgreSQL 16 (audit chain, queue, rules, users)
                        Redis 7 (rate limiting, transient state)
|
v
Dashboard               React 18 SPA, WebSocket push
```

Front end

React 18 with TypeScript, built by Vite 6.3. Tailwind CSS 4 for styling, Recharts for the dashboard visualisations, Radix primitives underneath the component set inherited from the A3 Figma work.

The SPA is a single bundle. `API_BASE` resolves to `import.meta.env.VITE_API_URL || ""`, so in the deployed configuration the frontend calls the same origin it was served from and needs no build-time API URL.

Live updates come over a WebSocket at `/api/v1/ws/live` rather than polling. The socket authenticates with `?token=<jwt>`; an unauthenticated socket is rejected. Polling exists only as a fallback if the socket cannot connect.

Back end

FastAPI on Python 3.13, served by uvicorn. Fourteen routers mounted under `/api/v1`: interception, audit, approval, policy, auth, stats, guardrails, websocket, agents, demo, reports, dryrun, trends and API keys.

Two pieces of middleware sit in front of every request: CORS, and a rate limiter allowing 60 requests/minute generally and 10/minute on auth endpoints. Behind a reverse proxy the limiter keys off `X-Real-IP`, because without it every client shares the proxy's bucket and one active dashboard locks everyone out.

Static files are mounted last, at `/`, so `/api/*` and `/health` keep winning against the SPA catch-all.

Interception path

`POST /api/v1/intercept` is the core of the system. Order of operations:

1. **Semantic guardrail runs first.** Regex pattern sets for prompt injection, jailbreak and exfiltration attempts, each carrying a hand-tuned confidence. A proximity-based discount suppresses the common false positive where text *discusses* injection rather than attempting it. If the guard blocks it short-circuits: outcome DENY, rule `SEMANTIC-GUARD`, risk Critical, and it still writes an audit entry. The policy engine never runs.
2. **Policy evaluation**, wrapped in try/except. Any exception becomes DENY at Critical risk.
3. **Approval item created** on ESCALATE only.
4. **Audit entry written for every outcome**, PERMIT included. No path through this endpoint skips the audit write.
5. **WebSocket broadcast** so connected dashboards update without polling.

Fail-closed behaviour

Four independent points, all real:

CONDITION	RESULT
No policy rule matches the action	DENY, risk Medium, "denied by default"
Policy engine raises	DENY, risk Critical
Policy YAML missing at startup	rules stay empty, so everything hits the no-match DENY
No valid API key or JWT	401, request never reaches policy

The system defaults to refusing rather than allowing.

Policy engine

Rules load from `backend/policies.yaml` (12 rules) via the `POLICY_CONFIG` setting. A further 10 pre-built rules ship in `rule_library.yaml` and can be imported into the database through the API, which is what populates the Policy Rules screen.

Conditions are parsed, not keyword-matched. The parser handles `==`, `!=`, `>`, `IN`, `NOT IN`, `MATCHES`, `NOT MATCHES` and `AND`, resolved against named sets (`approved_tools`, `approved_vendors`, `approved_domains`, `restricted_tables`, `restricted_datasets`) defined in the same file.

Where a condition cannot be resolved the rule **does not match**, and the request falls through to the default deny. That direction is deliberate: unresolvable conditions were previously treated as satisfied, which on a *permit* rule meant allowing exactly the thing the rule existed to gate.

Known limitation: `OR` is not supported.

Audit chain

Append-only and hash-chained. Each entry's hash is SHA-256 over a preimage built from every stored column: `entry_id`, `timestamp`, `agent_id`, `action_type`, `outcome`, `risk_level`, `policy_rule`, `policy_desc`, `session_id`, `payload_masked`, a `payload_hash`, and the previous entry's hash. The preimage is sorted JSON rather than concatenated strings, so two different entries cannot produce the same one by shifting where a field ends. The chain starts from a fixed genesis hash.

`payload_hash` is taken over the **raw** payload, before masking. That is deliberate - the integrity check then covers what the agent actually sent, not the redacted copy we display. The masked copy is covered separately as its own field.

Covering every column matters. The hash previously covered 7 of the 12 stored fields, so `risk_level` or the masked payload could be rewritten without breaking the chain: a Critical could be dropped to Low and verification

still reported valid. `GET /api/v1/audit/verify` now walks from genesis and recomputes each hash rather than only comparing stored hashes to one another.

Appends are serialised behind an `asyncio.Lock`. This is not incidental: two concurrent intercepts could previously both read the same chain head and both append to it, forking the log. `GET /api/v1/audit/verify` re-walks the chain server-side and recomputes every hash.

The guarantee is **tamper-evident, not tamper-proof**. Anyone with database access can still delete rows. What they cannot do is delete them without breaking the chain, and verification surfaces that. Append-only is enforced at application level; there is no database trigger.

Authentication and access control

JWT bearer tokens for users, bcrypt for password hashing. `auth.py` calls bcrypt directly and truncates to 72 bytes itself. That used to go through passlib, which probes its backend with an over-length string on import - bcrypt 4.1+ raises on that rather than truncating, so it broke every login on a clean install and was worked around by pinning `bcrypt==4.0.1`. Dropping passlib removed the pin along with the stale-pin problem; the hash format is unchanged, so existing rows verify.

Three roles: Administrator, Analyst, Viewer. Every endpoint carries an auth dependency; the role constants live in `backend/api/auth.py`. Agent-facing endpoints accept either an `X-API-Key` or a JWT via the `agent_or_user` dependency, because agents cannot perform the login flow.

Payloads are masked before storage. Masking covers top-level keys only, which is recorded as a limitation.

Data model

PostgreSQL 16, accessed asynchronously through SQLAlchemy 2.0 with asyncpg. Eight tables: `audit_log`, `approval_queue`, `policy_rules`, `users`, `agents`, `api_keys`, `scheduled_reports`, `report_exports`.

Entity relationship diagram, generated from the SQLAlchemy models in `backend/models/tables.py`: `docs/screenshots/database-erd.png` (mermaid source in `docs/database-erd.html`).

The relationships in that diagram are logical, not declared. No column carries a `ForeignKey` constraint, so the database will not reject an `audit_log` row naming an agent that was never registered. That is deliberate in one place and a known gap in the others:

- `audit_log.agent_id` **must** stay unconstrained. The whole point of the log is to record what an unknown or unregistered caller attempted. A foreign key here would make the system refuse to record exactly the events it exists to catch, and it would let anyone delete an agent row to orphan its history.
- `approval_queue.reviewer`, `policy_rules.created_by` and the `created_by` columns on `api_keys` and `scheduled_reports` hold an email rather than a user id, so a renamed account leaves the old string behind. These should be real references and are not.

Integrity is enforced in the application layer instead: the interception path resolves the agent before it writes counters, and the reviewer identity is taken from the verified JWT rather than the request body. Same register as the append-only guarantee below, which is also an application-level promise.

Alembic is configured, but the schema is created by `Base.metadata.create_all()` on startup and the one revision in `backend/alembic/versions` is empty. Nothing in the deployment runs `alembic upgrade`. Recorded here rather than left implied.

The agent registry was a dictionary in `backend/api/agents.py` until it moved into the `agents` table. Every restart of the container used to drop whatever had been registered through `POST /api/v1/agents` and put the hardcoded entry back.

Redis 7 backs rate limiting and transient state.

A background scheduler task starts with the application and drives scheduled report generation.

Deployment

Production is Docker Compose on an Azure VM, not a managed platform. Three services from `docker-compose.azure.yml`:

- `postgres` (postgres:16-alpine) with a named volume
- `redis` (redis:7-alpine)
- `app`, built from `infra/docker/Dockerfile.azure`, published `80:8000`

The app image is a two-stage build. Stage one runs `npm ci` and `vite build`; stage two installs the Python dependencies and copies the built frontend in as `frontend_dist/`, which FastAPI serves. One image, one container, no nginx.

Configuration is entirely environment variables (`DATABASE_URL`, `REDIS_URL`, `SECRET_KEY`, `AGENT_API_KEY`, `POLICY_CONFIG`, `RULE_LIBRARY`), read from a `.env` that is not in the repository. That makes the secrets externalisable to a secret manager without code changes.

Deployment is a pull-and-rebuild script on the VM, triggered over SSH by a GitHub Actions workflow. The CI deploy key is pinned to a forced command in the VM's `authorized_keys`, so it can run the deploy script and nothing else.

Non-functional characteristics

CONCERN	HOW IT IS ADDRESSED
Latency	WebSocket push rather than polling; sub-second dashboard updates
Integrity	SHA-256 hash-chained audit with server-side verification
Access control	JWT + bcrypt + RBAC on every endpoint
Confidentiality	PII masking on stored payloads (top-level keys)
Availability under load	Rate limiting, 60/min general and 10/min auth
Failure mode	Fail-closed at four independent points

Known gaps

Recorded here rather than only in the traceability matrix, because they are architectural rather than feature-level:

- Policy conditions do not support `OR` (FR-02).
- Approval SLA expiry is evaluated lazily on read; there is no background expiry worker (FR-05).
- PII masking covers top-level payload keys only.
- Audit append-only is enforced in the application, not by a database constraint.
- CORS origins are still the development list.