



GROUP 28 · CSIT321

Interface Design

Certacito.ai · A4 Prototype Presentation appendix

Project Certacito.ai, AI agent governance platform

Document interface-design.md

Date 07 August 2026

Live system <http://20.92.93.30>

A4 appendix. A3 presented eight high-fidelity Figma screens. A4 is those screens implemented as a working React app fed by the real API. This doc lists where the implementation deviates from the A3 designs and the reason for each change.

1. What stayed the same

The A3 design system survived contact with implementation almost untouched:

- Colour system and its "one job per colour" rule (navy structure, teal active/permit, red critical/deny, green permit/pass, amber escalate).
- Persistent navy sidebar, breadcrumb header, card-on-grey layout.
- Unified risk/outcome badges - colour always paired with a text label (WCAG 2.1 AA, risk is never conveyed by colour alone).
- Screen set and the two persona journeys (Compliance Officer triage path, DevSecOps configuration path).
- Expandable row detail on the audit log and approval queue.


2. What changed, and why

#	A3 DESIGN	A4 IMPLEMENTATION	WHY
1	Sign-in via Azure AD	Local JWT login (email + bcrypt password), SSO button kept as inactive placeholder	Live on an Azure VM in australiaeast. JWT + RBAC gives the same role-gated behaviour to demo; the SSO swap is isolated inside <code>auth.py</code> .
2	Typography: Inter / JetBrains Mono	Arial / Courier New system stack	Licensing-free system fonts render identically across marker machines with zero font loading. Visual intent (prose vs system identifiers) is preserved. Revisit for production.
3	Static mock numbers in every panel	All KPIs, feeds, charts and tables driven by the live API; WebSocket push for instant updates	The A3 prototype only had to look real. A4 has to be real - a live simulator (~4 events/min) plus the OpenClaw agent feed the same pipeline a customer deployment would.
4	Audit "Export modal"	Direct Export CSV (client-side) + Export PDF (print stylesheet) + a new Verify chain button	Verify chain surfaces the tamper-evidence story in one click during demos - it re-walks the whole SHA-256 chain server-side. The modal added a step without adding value.
5	No landing page in A3	Public marketing landing page before login	Exhibition feedback: visitors need the "what is this and why should I care" story before the console. Also carries the group branding.
6	Audit detail drawer showed one exemplar record	Detail row renders the actual selected entry: masked payload, its own chain hash, previous-entry hash, real decision trail	Markers click more than one row. Mock detail on real rows read as broken.
7	Notifications area (bell) with static badge	Badge counts live pending approvals + criticals	Same reason as 3.
8	2-second polling everywhere (A3 assumed it)	WebSocket-first, 10 s fallback poll	Polling three endpoints every 2 s tripped our own API rate limit after a minute on screen. The socket is instant and cheaper.


3. Screens as implemented

All ten screens of the implemented console, captured from the running Azure deployment with the simulator on (source files in docs/screenshots/). The A3 feedback asked to see every designed screen, so they are all here, ordered the way a user meets them: entry, then the Compliance Officer triage journey, then the DevSecOps configuration journey.

Landing page - the public "what is this" story before any login.




Login - JWT email/password with the SSO placeholder (change 1 in section 2).




Dashboard - live KPIs, intercepted-actions feed, risk mix. The Compliance

Officer's home screen.



Approval queue - the human-in-the-loop worklist; expandable rows carry the full decision context, approve/deny act on the live API.



Audit log - the hash-chained record with per-row chain hashes and the one-click Verify chain action.

certacito.ai

All systems operational

certacito.ai > Audit Log

LIVE
1
AD admin Administrator

MAIN

- Agent Overview
- Dashboard
- Activity**
- Audit log
- Approval queue

GOVERNANCE

- Policy rules
- Agent registry

REPORTING

- Reports & compliance

SETTINGS

- Settings
- Help & docs

admin

AD Administrator

CERTACITO.AI / ACTIVITY / AUDIT LOG

Agent Decision Audit Log

Review policy decisions, masked payloads, chain hashes and human escalation history.

Verify chain
Export CSV
Export PDF

TOTAL EVENTS 73 within current filtered view	DENIED 43 blocked by policy rule	ESCALATED 4 sent for human review	CRITICAL RISK 22 highest risk badge
---	---	--	--

All risk levels ▾
All outcomes ▾
All action types ▾
Reset

	TIMESTAMP	AGENT ID	ACTION TYPE	MATCHED POLICY RULE	RISK LEVEL	OUTCOME
▼	2026-08-07 08:08:14	AGT-gp-office-001	email_send	RULE-006 Matched RULE-006: External email send escalation	HIGH	ESCALATE

Full decision record

Audit event ID
Matched rule ID
Regulatory alignment
Policy decision

AUC-SA66C0D6
RULE-006
Internal policy
ESCALATE

```
to: specialist@external-clinic.com.au
subject: Patient Referral
has_pii: true
```

Payload has been masked before display to protect personal and secret data.

SHA-256 chain hash

```
c653d154725e8a072da46675a0637728e4aaa4e3fbf1d66dcf087bb9bdd3439
```

prev: cde7c7f69b263d1ad6e62e76f9190cf34ce8c353ba218e91ae8e08b1d43951c

Decision trail

- Policy engine matched RULE-006 - Matched RULE-006: External email send escalation
Step 1 · 2026-08-07 08:08:14
- Escalated to the human approval queue with an SLA deadline
Step 2 · 2026-08-07 08:08:14
- Awaiting or resolved by a reviewer - see approval queue
Step 3

Reports and compliance - exportable compliance posture over time.

Agent overview - fleet status at a glance; the DevSecOps entry point.

certacito.ai

All systems operational

MAIN

Agent Overview

Dashboard

Activity

GOVERNANCE

Policy rules

Agent registry

REPORTING

Reports & compliance

Settings

Help & docs

Policy engine online · 22ms

admin Administrator

certacito.ai

Agent Overview

Search agents, rules, audit events...

LIVE

admin Administrator

CERTACITO.AI / AGENT OVERVIEW

Agent Overview

All registered AI agents and their current governance status.

Last updated: just now

LIVE

1 Total Agents

1 Active

0 Suspended

0 Offline

AGT-opencLaw-azure

OpenClaw Governance Agent

Active Low

Compliance 100%

View in Registry

Agent registry - per-agent detail, permitted action types, live counters.

certacito.ai

All systems operational

MAIN

Agent Overview

Dashboard

Activity

GOVERNANCE

Policy rules

Agent registry

REPORTING

Reports & compliance

Settings

Help & docs

Policy engine online · 22ms

admin Administrator

certacito.ai

Agent Registry

Search agents, rules, audit events...

LIVE

admin Administrator

CERTACITO.AI / AGENT REGISTRY

Agent Registry

Manage and monitor all registered AI agents across your organisation

Refresh Register agent

TOTAL AGENTS 1

ACTIVE 1

SUSPENDED 0

HIGH / CRITICAL RISK 0

DECISIONS TODAY 0

Search by agent ID, name, type, or team...

All statuses

All risk tiers

All environments

AGENT ID	NAME / TYPE	STATUS	RISK	DECISIONS	VIOLATIONS	COMPLIANCE	LAST ACTIVE	ACTIONS
AGT-opencLaw-azure	OpenClaw Governance Agent google/gemini-2.5-pro	Active	LOW	0 today	0	100%	no activity yet	<button>View</button> <button>Suspend</button>

Showing 1 of 1 agents

Production: 1 Staging: 0

Policy rules - policy CRUD plus the pre-built rule library.

CERTACITO.AI

University of Wollongong | CSIT321 | 2026 | Group 28

certacito.ai

All systems operational

MAIN

Agent Overview

Dashboard

Activity

Audit log

Approval queue

GOVERNANCE

Policy rules

Agent registry

REPORTING

Reports & compliance

Settings

Help & docs

Policy engine online · 22ms

AD admin Administrator

certacito.ai > Policy Rules

Search agents, rules, audit events...

LIVE

AD admin Administrator

CERTACITO.AI / POLICY RULES

Policy Rules

Manage enforcement rules for AI agent governance

+ Create new rule

Rules are stored as version-controlled config files, not model prompts. All changes are tracked with git-style versioning.

All

prompt_content

*

data_access

external_call

credential_access

file_write

Rule ID	Name	Action type	Risk threshold	Default outcome	Regulatory tag	Status	Last modified
LIB-AI-001	Block prompt injection in t...	prompt_content	CRITICAL	DENY	OWASP LLM01	Active	2026-08-05 v1
LIB-AI-002	Block sensitive data in LL...	prompt_content	HIGH	DENY	OWASP LLM06 + AU Guardrail 7	Active	2026-08-05 v1
LIB-AI-003	Require human approval f...	*	HIGH	ESCALATE	AU Guardrail 3 + AU Guardrail 6	Active	2026-08-05 v1
LIB-FIN-001	Block bulk financial data e...	data_access	HIGH	DENY	APRA CPS 234	Active	2026-08-05 v1
LIB-FIN-002	External transfer requires ...	external_call	CRITICAL	ESCALATE	APRA CPS 234 + AUSTRAC	Active	2026-08-05 v1
LIB-HEALTH-001	Patient record access req...	data_access	CRITICAL	DENY	Privacy Act 1988 s13G	Active	2026-08-05 v1
LIB-					My Health		

Settings - organisation, team, notifications, API keys, integrations. Team members and API keys are live (real users from the API, keys issued and revoked from this screen); the organisation form and notification toggles are client-side only in this release, and the integrations tab honestly reports every integration as not connected rather than faking a status.

certacito.ai

All systems operational

MAIN

Agent Overview

Dashboard

Activity

Audit log

Approval queue

GOVERNANCE

Policy rules

Agent registry

REPORTING

Reports & compliance

Settings

Help & docs

Policy engine online · 25ms

AD admin Administrator

certacito.ai > Settings

Search agents, rules, audit events...

LIVE

AD admin Administrator

CERTACITO.AI / SETTINGS

Settings

Organisation configuration, team management, and platform integrations

Organisation

Team Members

Notifications

API Keys

Integrations

AGENT LIMIT

180 max

Set below under Governance Limits

Governance Limits

Maximum Registered Agents

180 agents

Audit Log Retention

90 days

Discard changes

Save changes

4. Sponsor/supervisor feedback applied

- "Interception-first framing" (Anthony, A3 review) - the dashboard leads with the

live intercepted-actions feed, not agent inventory.

- "Understandable in a few minutes by someone new" (supervisor email, June) - the landing page + healthcare demo scenario (POST /api/v1/demo/healthcare-scenario) give a self-contained walkthrough.

5. A3 marking feedback applied

The A3 assessment feedback listed specific improvements. Point by point:

A3 FEEDBACK	WHAT WE DID IN A4
"Only a limited number of interface screens are presented"	All ten screens are implemented and shown in section 3, captured from the live system - none are mockups
"Additional screenshots for all designed screens would make the screen flow even clearer"	The section 3 gallery is ordered along the two persona journeys, so the flow reads top to bottom
Design approval "should be presented more consistently"	The approval record now lives in exactly one place - the sign-off box in the contribution table appendix - and no other document makes an approval claim
"Provide additional user testing or usability evaluation"	New Usability Evaluation appendix: a heuristic walkthrough and WCAG 2.1 AA audit of the live deployment, plus the sponsor, supervisor and exhibition feedback and what each changed
"Continue implementing the frontend and backend to demonstrate the complete governance platform"	The A4 submission is that: the full platform live on Azure, 105 passing tests, and a real agent (OpenClaw) governed through it
"Ensure all group members actively contribute during demonstrations and clearly explain their individual design contributions"	Each member's individual contribution is recorded and signed in the contribution table appendix