



GROUP 28 · CSIT321

Interface Design

Certacito.ai · A4 Prototype Presentation appendix

Project Certacito.ai, AI agent governance platform

Document interface-design.md

Date 07 August 2026

Live system <http://20.92.93.30>

A4 appendix. A3 presented eight high-fidelity Figma screens. A4 is those screens implemented as a working React app fed by the real API. This doc lists where the implementation deviates from the A3 designs and the reason for each change.

1. What stayed the same

The A3 design system survived contact with implementation almost untouched:

- Colour system and its "one job per colour" rule (navy structure, teal active/permit, red critical/deny, green permit/pass, amber escalate).
- Persistent navy sidebar, breadcrumb header, card-on-grey layout.
- Unified risk/outcome badges - colour always paired with a text label (WCAG 2.1 AA, risk is never conveyed by colour alone).
- Screen set and the two persona journeys (Compliance Officer triage path, DevSecOps configuration path).
- Expandable row detail on the audit log and approval queue.

2. What changed, and why

#	A3 DESIGN	A4 IMPLEMENTATION	WHY
1	Sign-in via Azure AD	Local JWT login (email + bcrypt password), SSO button kept as inactive placeholder	Live on an Azure VM in australiaeast. JWT + RBAC gives the same role-gated behaviour to demo; the SSO swap is isolated inside <code>auth.py</code> .
2	Typography: Inter / JetBrains Mono	Arial / Courier New system stack	Licensing-free system fonts render identically across marker machines with zero font loading. Visual intent (prose vs system identifiers) is preserved. Revisit for production.
3	Static mock numbers in every panel	All KPIs, feeds, charts and tables driven by the live API; WebSocket push for instant updates	The A3 prototype only had to look real. A4 has to be real - a live simulator (~4 events/min) plus the OpenClaw agent feed the same pipeline a customer deployment would.
4	Audit "Export modal"	Direct Export CSV (client-side) + Export PDF (print stylesheet) + a new Verify chain button	Verify chain surfaces the tamper-evidence story in one click during demos - it re-walks the whole SHA-256 chain server-side. The modal added a step without adding value.
5	No landing page in A3	Public marketing landing page before login	Exhibition feedback: visitors need the "what is this and why should I care" story before the console. Also carries the group branding.
6	Audit detail drawer showed one exemplar record	Detail row renders the actual selected entry: masked payload, its own chain hash, previous-entry hash, real decision trail	Markers click more than one row. Mock detail on real rows read as broken.
7	Notifications area (bell) with static badge	Badge counts live pending approvals + criticals	Same reason as 3.
8	2-second polling everywhere (A3 assumed it)	WebSocket-first, 10 s fallback poll	Polling three endpoints every 2 s tripped our own API rate limit after a minute on screen. The socket is instant and cheaper.

3. Screens as implemented

Current screenshots live in `docs/screenshots/` (taken from the running system, simulator on). Screen-by-screen purpose is unchanged from A3 section 3.3.

4. Sponsor/supervisor feedback applied

- "Interception-first framing" (Anthony, A3 review) - the dashboard leads with the live intercepted-actions feed, not agent inventory.
- "Understandable in a few minutes by someone new" (supervisor email, June) - the landing page + healthcare demo scenario (`POST /api/v1/demo/healthcare-scenario`) give a self-contained walkthrough.