



GROUP 28 · CSIT321

# Interface Design

---

Certacito.ai · A4 Prototype Presentation appendix

**Project** Certacito.ai, AI agent governance platform

**Document** interface-design.md

**Date** 07 August 2026

**Live system** <http://20.92.93.30>

A4 appendix. A3 presented eight high-fidelity Figma screens. A4 is those screens implemented as a working React app fed by the real API. This doc lists where the implementation deviates from the A3 designs and the reason for each change.

## 1. What stayed the same

The A3 design system survived contact with implementation almost untouched:

- Colour system and its "one job per colour" rule (navy structure, teal active/permit, red critical/deny, green permit/pass, amber escalate).
- Persistent navy sidebar, breadcrumb header, card-on-grey layout.
- Unified risk/outcome badges - colour always paired with a text label (WCAG 2.1 AA, risk is never conveyed by colour alone).
- Screen set and the two persona journeys (Compliance Officer triage path, DevSecOps configuration path).
- Expandable row detail on the audit log and approval queue.

## 2. What changed, and why

| # | A3 DESIGN                                      | A4 IMPLEMENTATION                                                                                                          | WHY                                                                                                                                                                             |
|---|------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | Sign-in via Azure AD                           | Local JWT login (email + bcrypt password), SSO button kept as inactive placeholder                                         | Live on an Azure VM in australiaeast. JWT + RBAC gives the same role-gated behaviour to demo; the SSO swap is isolated inside <code>auth.py</code> .                            |
| 2 | Typography: Inter / JetBrains Mono             | Arial / Courier New system stack                                                                                           | Licensing-free system fonts render identically across marker machines with zero font loading. Visual intent (prose vs system identifiers) is preserved. Revisit for production. |
| 3 | Static mock numbers in every panel             | All KPIs, feeds, charts and tables driven by the live API; WebSocket push for instant updates                              | The A3 prototype only had to look real. A4 has to be real - a live simulator (~4 events/min) plus the OpenClaw agent feed the same pipeline a customer deployment would.        |
| 4 | Audit "Export modal"                           | Direct Export CSV (client-side) + Export PDF (print stylesheet) + a new <b>Verify chain</b> button                         | Verify chain surfaces the tamper-evidence story in one click during demos - it re-walks the whole SHA-256 chain server-side. The modal added a step without adding value.       |
| 5 | No landing page in A3                          | Public marketing landing page before login                                                                                 | Exhibition feedback: visitors need the "what is this and why should I care" story before the console. Also carries the group branding.                                          |
| 6 | Audit detail drawer showed one exemplar record | Detail row renders the actual selected entry: masked payload, its own chain hash, previous-entry hash, real decision trail | Markers click more than one row. Mock detail on real rows read as broken.                                                                                                       |
| 7 | Notifications area (bell) with static badge    | Badge counts live pending approvals + criticals                                                                            | Same reason as 3.                                                                                                                                                               |
| 8 | 2-second polling everywhere (A3 assumed it)    | WebSocket-first, 10 s fallback poll                                                                                        | Polling three endpoints every 2 s tripped our own API rate limit after a minute on screen. The socket is instant and cheaper.                                                   |

### 3. Screens as implemented

Current screenshots live in [docs/screenshots/](#) (taken from the running system, simulator on). Screen-by-screen purpose is unchanged from A3 section 3.3.

### 4. Sponsor/supervisor feedback applied

- "Interception-first framing" (Anthony, A3 review) - the dashboard leads with the live intercepted-actions feed, not agent inventory.
- "Understandable in a few minutes by someone new" (supervisor email, June) - the landing page + healthcare demo scenario ( [POST /api/v1/demo/healthcare-scenario](#) ) give a self-contained walkthrough.